1   23rd of November 2018 — F. Poloni

1.1   Gaussian elimination on symmetric matrices

> **💡 Do you recall?**
>
> In Gaussian elimination we had A and we multiplied it by L_1 in order to get a big chunk of 0s in the first column
>
> $$\begin{pmatrix} 1 & & & & \\ * & 1 & & & \\ * & & 1 & & \\ * & & & 1 & \\ * & & & & 1 \end{pmatrix}\begin{pmatrix} \circledast & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} = \begin{pmatrix} \circledast & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \end{pmatrix}$$

Let us consider an upgrade of Gaussian elimination in the case of $A \in S(m, \mathbb{R})$.

Let us see what happens if we multiply $L_1 A$ on the right by the transpose of L_1:

STEP 1:

$$\begin{pmatrix} 1 & & & & \\ * & 1 & & & \\ * & & 1 & & \\ * & & & 1 & \\ * & & & & 1 \end{pmatrix}\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}\begin{pmatrix} 1 & * & * & * & * \\ & 1 & & & \\ & & 1 & & \\ & & & 1 & \\ & & & & 1 \end{pmatrix} = \begin{pmatrix} * & 0 & 0 & 0 & 0 \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \end{pmatrix}$$

STEP 2:

$$\begin{pmatrix} 1 & & & & \\ & 1 & & & \\ & * & 1 & & \\ & * & & 1 & \\ & * & & & 1 \end{pmatrix}\begin{pmatrix} * & 0 & 0 & 0 & 0 \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \end{pmatrix}\begin{pmatrix} 1 & & & & \\ & 1 & * & * & * \\ & & 1 & & \\ & & & 1 & \\ & & & & 1 \end{pmatrix} = \begin{pmatrix} * & 0 & 0 & 0 & 0 \\ 0 & * & 0 & 0 & 0 \\ 0 & 0 & * & * & * \\ 0 & 0 & * & * & * \\ 0 & 0 & * & * & * \end{pmatrix}$$

STEP m:

$$L_{m-1}L_{m-2}\ldots L_1 A L_1^T \ldots L_{m-2}^T L_{m-1}^T = D$$

where D is diagonal, or:

$$A = L_1 L_2 \ldots L_{m-1} D L_{m-1}^T \ldots L_2^T L_1^T = LDL^T$$

Observation 1.1 (Stroke of luck)**.** *Notice that the stroke of luck of **??** holds in this case too, hence we pay nothing to compute matrix L.*

$$\begin{bmatrix} 1 & & & & \\ -a_2 & 1 & & & \\ -a_3 & & 1 & & \\ -a_4 & & & 1 & \\ -a_5 & & & & 1 \end{bmatrix}^{-1}\begin{bmatrix} 1 & & & & \\ & 1 & & & \\ & -b_3 & 1 & & \\ & -b_4 & & 1 & \\ & -b_5 & & & 1 \end{bmatrix}^{-1}\begin{bmatrix} 1 & & & & \\ & 1 & & & \\ & & 1 & & \\ & & -c_4 & 1 & \\ & & -c_5 & & 1 \end{bmatrix}^{-1}\begin{bmatrix} 1 & & & & \\ & 1 & & & \\ & & 1 & & \\ & & & 1 & \\ & & & -d_5 & 1 \end{bmatrix}^{-1} = \begin{bmatrix} 1 & & & & \\ a_2 & 1 & & & \\ a_3 & b_3 & 1 & & \\ a_4 & b_4 & c_4 & 1 & \\ a_5 & b_5 & c_5 & d_5 & 1 \end{bmatrix}$$

Theorem 1.1 (Symmetric Gaussian elimination)**.** *Let $A \in S(m, \mathbb{R})$ such that during Gaussian elimination we don't encounter any 0 pivot. A admits a factorization $A = LDL^T$, where L is lower triangular with ones on its diagonal, and D is diagonal.*

A Matlab implementation of symmetric Gaussian elimination is shown in Algorithm 1.1.

ALGORITHM 1.1 Symmetric Gaussian factorization, Matlab implementation.

```matlab
function [L, D] = ldl_factorization(A)
  m = size(A, 1);
  L = eye(m); D = zeros(m);
  for k = 1:m-1
    D(k, k) = A(k, k);
        L(k+1:end, k) = A(k+1:end, k) / A(k, k);
        A(k+1:end, k+1:end) = A(k+1:end, k+1:end) ...
      - L(k+1:end, k) * A(k, k+1:end);
  end
  D(m, m) = A(m, m);
```

It is possible to make an optimization of this algorithm: since A is supposed to be symmetric, we only need to update the lower triangular part of A, since the rest is mirrored by symmetry, hence the computational complexity is half that of Gaussian elimination.

This algorithm is not backward stable, exactly like the one on non symmetric matrices. Pivoting may be performed in order to improve stability. It comes without saying that the row swap should be done consistently on the columns to preserve symmetry.

Of course there are some matrices (like the ones with all 0s on the diagonal) that cannot be "pivoted". There are workarounds, though. As an example, Matlab's `[L, D, P] = ldl(A)` produces matrices such that $P^T A P = LDL^T$, where D may have 2×2 diagonal blocks.

> 💡 **Do you recall?**
>
> We recall the characterization of **positive definite matrix** $A \in M(m, \mathbb{R})$: all its eigenvalues are strictly positive. In other words, A is positive definite if $\forall z \neq 0 \in \mathbb{R}^m \; z^T A z > 0$.

Lemma 1.2. *In the context of positive definite matrices the following holds:*

1. *Let A be a symmetric matrix. A is positive definite if and only if MAM^T is so, for some invertible $M \in M(m, \mathbb{R})$. Formally, $\forall A \in S(m, \mathbb{R})$ s.t. $A \succ 0 \Leftrightarrow \exists M \in GL(m, \mathbb{R})$ s.t. $MAM^T \succ 0$;*

2. Let A a symmetric positive definite matrix such that $A = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}$, then A_{11} and A_{22} are, too. Formally, $\forall A \in S(m, \mathbb{R})$ s.t. $A \succ 0$ and $A = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \Rightarrow A_{11} \succ 0$ and $A_{22} \succ 0$.

Proof.

1. $\Rightarrow$) $A \in S(m, \mathbb{R})$ and $A \succ 0 \Longrightarrow MAM^T \in S(m, \mathbb{R})$ and $MAM^T \succ 0$.

   Take $z \in \mathbb{R}^m$, $z \neq 0$ $z^T MAM^T z = y^T Ay > 0$, where we performed a variable change $y = M^T z$. Notice that $y \neq 0$ because M is invertible (and $ker(M) = \{0\}$). The symmetry of the matrix MAM^T follows from $(MAM^T)^T = M^{T^T} A^T M^T = MAM^T$;

   $\Leftarrow$) $MAM^T \in S(m, \mathbb{R})$ and $MAM^T \succ 0 \Longrightarrow A \in S(m, \mathbb{R})$ and $A \succ 0$.

   This proof follows from the previous arrow, where the substitution is $z = M^{-1}y$.

2. $A = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}$ positive definite $\Longrightarrow A_{11}$ and A_{22} are positive definite too.

   Since A is positive definite, its scalar product is greater than zero with all the vectors in $\mathbb{R}^m$.

   A_{11}) Let us take $\mathbf{z} = \begin{bmatrix} \mathbf{z_1} \\ 0 \end{bmatrix}$

   $\begin{bmatrix} \mathbf{z_1}^T & 0 \end{bmatrix} \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} \mathbf{z_1} \\ 0 \end{bmatrix} = \mathbf{z_1}^T A_{11} \mathbf{z_1} > 0, \forall \mathbf{z_1} \in \mathbb{R}^{sizeof A_{11}}$

   A_{22}) Let us take $\mathbf{z} = \begin{bmatrix} 0 \\ \mathbf{z_2} \end{bmatrix}$

   $\begin{bmatrix} 0 & \mathbf{z_2}^T \end{bmatrix} \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} 0 \\ \mathbf{z_2} \end{bmatrix} = \mathbf{z_2}^T A_{22} \mathbf{z_2} > 0, \forall \mathbf{z_2} \in \mathbb{R}^{sizeof A_{22}}$

$\square$

Corollary 1.3. *Let $A \in M(n, \mathbb{R})$ such that A is positive definite. When computing the LDL^T factorization of A, at each step we have $D_{kk} > 0$, hence we need no pivoting technique.*

Proof. From the first point of Lemma 1.2 we have that, since $A \succ 0$, $L_1 A L_1^T$ is positive definite. Thanks to the second point of the same lemma we have $L_1 A L_1^T = \begin{pmatrix} * & 0 & 0 & 0 & 0 \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \\ 0 & * & * & * & * \end{pmatrix}$

and so the first and the second diagonal blocks are positive definite ($D_{11} \succ 0$ and $D_{22} \succ 0$).

Notice that, since A is positive definite $A_{11} \succ 0$, but it's a scalar, hence $A_{11} > 0$ and this implies no breakdown case. $\square$

We may introduce another kind of factorization.

1.2 Cholesky factorization

The key idea is to write the diagonal matrix of the Gaussian elimination D as product of $D^{1/2}$ times itself:

$$D = \begin{pmatrix} d_{11} & & & \\ & d_{22} & & \\ & & \ddots & \\ & & & d_{mm} \end{pmatrix} = \begin{pmatrix} \sqrt{d_{11}} & & & \\ & \sqrt{d_{22}} & & \\ & & \ddots & \\ & & & \sqrt{d_{mm}} \end{pmatrix} \begin{pmatrix} \sqrt{d_{11}} & & & \\ & \sqrt{d_{22}} & & \\ & & \ddots & \\ & & & \sqrt{d_{mm}} \end{pmatrix}$$

The LDL factorization may be rewritten as follows:

$$A = LDL^T = LD^{1/2}(D^{1/2^T}L^T) = CC^T$$

where $D^{1/2} = \text{diag}(D_{11}^{1/2}, D_{22}^{1/2}, \ldots, D_{mm}^{1/2})$, and C is lower triangular (but not anymore with ones on the diagonal).

In Matlab the Cholesky factorization of a positive definite matrix is performed by the function `chol(A);` and returns C^T.

Observation 1.2. *We will not discuss stability further, but Cholesky is always backward stable even without pivoting ($\|C\| = \|A\|^{1/2}$).*

Observation 1.3. *In a sparse matrix, we can choose the (symmetric) permutation with the only goal of reducing fill-in. The same considerations about LU factorization hold in this case too.*

1.3 Krylov subspace methods

In this part we will discuss some different techniques to solved linear systems, inspired from optimization algorithms.

Example 1.1. *Let us consider the optimization problem* $\min \frac{1}{2}x^T A x + b^T x$, *where $A \succ 0$. We know that the solution to this problem is $x = A^{-1}b$. Assume we solve this problem via a gradient descent-type method, starting from $x_0 = 0$.*

STEP 1: *$x_0 = 0$, $\nabla f(x_0) = -b$;*

STEP 2: *$x_1 =$ some multiple of $b \in Span(b)$*

$f(x_1) = Ax_1 - b$

$\nabla f(x_1) =$ *some multpiple of Ab + some multiple of $b \in span(b)$;*

STEP 3: *$x_2 =$ mult. of x_1 + mul. of $\nabla f(x_1)$ + mult. of $x_0 + \cdots \in span(Ab_1, b)$;*

STEP 4: *$x_3 =$ mult. of $x_2 + \cdots \nabla f(x_2)$+ previous iterates $\cdots$*

$Ax_2 - b = A(sAb + tb) - b = sA^2b + tAb - b \in span(b, Ab, A^2b)$, *where s and t are scalars;*

STEP 5: $x_4 \in span(b, Ab, A^2b, A^3b)$.

Notice that we can make some linear combinations of the vectors we have available and $A(A(sb + tAb) + uAb + vb) + e(sb + tAb) + fAb + gb \in span(b, Ab, A^2b, A^3b)$.

Idea: first compute the basis of $span(b, Ab, A^2b, A^3b)$, then look for the best solution inside this subspace.

The following family of algorithms "uses" the matrix A only by computing matrix-vector products.

Observation 1.4. *The cost of multiplying a sparse matrix A with a vector z is $O(nnz(A))$, where $nnz(A)$ is the number of non zero entries of A.*

Proof. Let us assume the matrix A is stored as a vector, whose entries are (i, j, A_{ij}).
The matrix-vector product would then be

1. x = zeros

2. for (i, j, A_{ij}) such that $A_{ij} \neq 0$

3. $x_i = x_i + A_{ij} * z_j$

$\square$

From now on we will consider to have a function `compute_product_with_A` that we use to compute matrix-vector products with A. In particular, $x =$`compute_product_with_A(z)` computes $x = Az$, given z. This function will be the only way in which the matrix A appears in our algorithm. If A is sparse, hence, these algorithms become particularly fast. Moreover, if we somehow have matrices that are not really sparse, but for which there exists a clever implementation of the matrix-vector product, this class of algorithms will give good results.

Definition 1.1 (Krylov subspace). *Let $A \in M(m, \mathbb{R})$ and let $b \in \mathbb{R}^m$. The **Krilov subspace** of index n is $\mathbf{K_n}(\mathbf{A}, \mathbf{b}) = span(b, Ab, A^2b, \ldots, A^{n-1}b)$.*
Equivalently, $k \in K_n(A, b) \iff \exists \alpha_1, \ldots, \alpha_{n-1} \in \mathbb{R}^m$ s.t. $v = \alpha_0 b + \alpha_1 Ab + \alpha_2 A^2b + \cdots + \alpha_{n-1} A^{n-1}b$.
Equivalently, $(\alpha_0 + \alpha_1 A + \alpha_2 A^2 + \cdots + \alpha_{n-1} A^{n-1})b = p(A)b$ for a polinomial p of degree such that $deg(p) \leq n - 1$.

Observation 1.5 (Properties).

1. *$v, w \in K_n(A, b) \Rightarrow \alpha + \beta W \in K_n(A, b)$;*

2. *$v \in K_n(A, b) \Rightarrow Av \in K_n(A, b)$;*

   Proof.
   Let us take $v = \alpha_0 b + \cdots + \alpha_{n-1} b$, then $Av = A(\alpha_0 b + \cdots + \alpha_{n-1} b) = \alpha_0 Ab + \cdots + \alpha_{n-1} A^n b$. $\square$

3. $\dim(K_n(A, b)) \leq n$. It is exactly n if $b, Ab, A^2b, \ldots, A^{n-1}b$ are linearly independent;

4. let us assume $dim(K_n(A, b)) \leq n$. In the second point, if A^{n-1} was really necessary $\alpha_{n-1} \neq 0$ or $v \in K_n(A, b)$, $v \notin K_{n-1}(A, b)$, equivalently then $A^n b$ is really necessary to write Av, i.e. $Av \in K_{n+1}(A, b)$ but $Av \notin K_n(A, b)$;

5. we may observe that $\dim(K_1(A, b)) < \dim(K_2(A, b)) < \cdots < \dim(K_{n_{max}}(A, b)) = \dim(K_{n_{max}+1}(A, b)) = \cdots$.